4

Comparing Participation Using Qualitative Comparative Analysis

Why now for a cumulative qualitative comparison?

Designing social research involves trade-offs between complexity and generality. There remains an often overblown and unhelpful distinction in social sciences between strategies that standardly prioritize one over the other. For Ragin, 'it is easy to exaggerate their differences and to caricature the two approaches, for example, by portraying quantitative work on general patterns as scientific but sterile and oppressive and qualitative research on small Ns as rich and emancipatory but soft and subjective' (2000: 22). QCA sets itself up as a method that provides an avenue for assuaging some of the unhelpful tensions in this distinction.

The attractiveness of QCA as both a critical methodological approach and a novel set of techniques has led to its application across an ever-expanding group of diverse disciplines and subdisciplines in the social sciences (Thiem and Dusa, 2013: 2). It has been one of the few approaches in recent times to come from outside the methodological orthodoxy that has not been quickly chased back to the margins of social science. The number of peer-reviewed articles per year using QCA has grown gradually since the first in 1984 – up to 15 in 2005, to 35 in 2011, and then more quickly to 45 in 2012 and 99 in 2013.[1] Of those 99 articles, 16 were agenda-setting articles introducing QCA to new subdisciplines (among them, Ryan and Smith, 2012). A large and yet non-exhaustive bibliography of QCA studies on the

[1] COMPASSS Newsletter, Nr 33, 10 January 2014. Available at: https://compasss.org/wp-content/uploads/simple-file-list/compasss-newsletter-33-Jan2014.pdf

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

COMPASSS[2] website attests to the method's continued growth in popularity, containing over 1,000 articles and 80 books at the time of writing in 2020.

There is a well-worn debate in the social sciences about the value of large-N and small-N research strategies.[3] Part of the early promise of QCA was to develop a method that makes feasible medium-N research strategies, in particular by extending the logic of small-N comparisons in a reasonable way. Ragin (2000: 25) plots the relative number of studies against the number of cases per study in the discipline producing a U-shaped curve, with several studies covering a handful of cases, several covering hundreds of cases, and a few covering case numbers that lie in between. The finding suggests that researchers are indeed quite averse to increasing the number of observations unless the number of cases can be radically improved (which might often require a significant change of research question).

Despite advances in large-N research applications, to answer many of the questions students of social phenomena would like to answer a large-N is notoriously hard to come by (Lijphart, 1971), or perhaps actually undesirable (Flyvbjerg, 2006). This can be particularly true of public policy innovation, as in the case of PB where diffusion is common but rarely universal, and intricate exemplars are valued by end-users of research. Many publications on PB have borne witness to its wide dispersal. There are in-depth accounts of celebrated exemplars, but scholars have struggled to find a unifying framework to compare cases without losing the rich substantive knowledge of context that seems so important to explaining what happened in those cases. Small-N case-based research can be attentive to the complexity of phenomena, allowing researchers to reflectively refine definitions of the elemental qualities of the things they are studying within the process of undertaking research. The trade-off is a lack of generalizability that decreases the potential for social research to be able to justify action and make an impact.

The withholding of generalization as a short-term goal is often the gambit of more constructivist approaches to research. A little bit of the generalizing discipline of existing conceptions is forsaken in the beginning, in order that it can be reintroduced later. The danger of this approach is that in allowing creative freedom, concepts become stretched

[2] COMPASSS (COMPArative Methods for Systematic cross-caSe analySis) is a worldwide network for scholars interested in systematic cross-case comparative approaches to research: https://compasss.org

[3] For a more in-depth treatment of my thoughts on this debate, see Ryan (2016, 2017, 2018).

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

beyond useful meaning. Yet, when done well, constructivist approaches can lead to accurate refinement of ideas to reflect observations.

Social science continues to divide between research that lumps the cases together while studying a handful of things that vary across them, and research that lumps the variables together and studies their intricate connections across a handful of cases. Few people venture in between. This alone is not a good enough reason to do so – people avoid risky journeys for very good reasons. However, Berg-Schlosser and Cronqvist suggest that QCA may offer a vehicle for the venture:

> Between the extremes of over-generalizing and 'universalizing' macro-quantitative approaches, on the one hand, and purely individualizing case-oriented approaches, on the other, a meaningful 'medium-range' social science can be built which, at the same time, has a higher explanatory power and a greater social and political relevance. (2005: 172)

PB is ready for a medium-N study of what PB is and what it does. QCA is a method that can grasp the opportunity provided by the build-up of worldwide case knowledge and theory but overcome the complexity and disparate nature of the data and the theories it mobilizes. It provides a set of tools to improve the interpretation of several cases where the case researchers may otherwise have their eyes close to the data such that it is difficult to zoom out and take a systematic approach to the evidence. Using QCA we can test the types of causal claims and intuitions of the researchers whose case studies and small-N comparisons have paved the way in studying democratic innovations. The following sections further elaborate the distinctiveness of the approach, and apply the lessons of recent methodological debates to substantive investigations of democratic innovation.

What is qualitative comparative analysis?

Despite its growth in popularity, many readers may still be new to the method and approach of QCA, which comes with its own set of terms and assumptions. Introducing vast new terminology to their audience is a dilemma for those writing about fuzzy sets.[4] I try to overcome that burden for the unfamiliar reader by introducing and returning to

[4] For book-length treatments of fsQCA methodology, see Ragin (2000, 2008), Smithson and Verkuilen (2006), Rihoux and Ragin (2007) and Schneider and Wagemann (2012).

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

the concepts gradually throughout the examples and analysis in the remainder of the book, to highlight the advantages of the method for knowledge accumulation. Nevertheless, in this section I briefly outline some key concepts and terms as a first point of reference for readers less familiar with the method, showing where relevant how they correspond to theories in democratic innovation research.

Necessity and sufficiency

QCA is an excellent method for testing the combinations of a number of conditions that are associated by set-theoretic relationships with an outcome. The logic underlying the way we try to understand the social world is often set-theoretic (see Schneider and Wagemann, 2012: 5–6). QCA explores set-theoretic relationships among variables or conditions such as relations of *necessity* and *sufficiency* (Schneider and Wagemann, 2012). The underlying logic is aimed at testing for 'If X, then Y' relationships rather than 'The more of X, the more of Y' relationships among conditions (Thiem et al, 2016).

Theorists and case researchers are particularly fond of claims of the nature 'X is sufficient but not necessary for Y'. Sufficient conditions (or more likely sufficient conjunctions of conditions) are often what many of us have in mind when we say out loud, 'X leads to Y'. We often use imperfect synonyms, which can give conditional or causal interpretations to necessity and sufficiency – such as X is 'required' or 'enough' for Y.

Take the following abbreviated passage from Gret and Sintomer's book on PB in Porto Alegre,

> A combination of a strong political culture in civil society and a strong political will in government represents the best context for initiating such an experiment. The process also requires a level of pragmatism in its implementation … formal political equality is not sufficient … true political will is required founded upon alternative action mechanisms … municipal action alone cannot hope to modify the relationship between social classes. (2005: 134)

In the space of half a page the authors suggest at least four hypotheses that are set-theoretic in nature. The hypotheses they are conceiving are based on notions of necessary (or 'required') and sufficient or insufficient ('not alone') conditions. The use of set-theoretic reasoning is ubiquitous in the social sciences, and statements about necessity

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

and sufficiency are rife within the PB literature. These kinds of relationships among conditions can be identified analytically in *subset–superset* relations among explanatory conditions (often called *influencing conditions*) compared across cases. We can establish necessity when we observe a set of cases containing the outcome to be a subset of the set of cases displaying the cause. Similarly, for sufficiency to be established, the set of cases containing the causal condition must be a subset of the cases displaying the outcome (Ragin, 2000: 214–17).[5]

Combinatorial and asymmetric causation and explanation

These hypotheses are also based on combinatorial or *conjunctural* causation. Configurative methods like QCA ascribe to an understanding of causality in the social world in which outcomes are often best explained by conjunctions of conditions. An analogy with elemental chemistry is useful where social conditions are like elements that can be combined in different ways to fashion more complex and in some ways unrecognizable compounds with very different properties. In configurative approaches to understanding cases in social science, change in the presence or absence of a single condition is understood to signify a different kind of case. It is the mix of conditions in a case that matters.

QCA provides a system of analysis that retains this complexity of cases, and yet uncovers and describes evidence for multiple conjunctural causation. The implication is that we can find relatively concise explanations across cases that still allow for *equifinality* – where the same outcome can be explained by different combinations of conditions. This equifinality assumption allows for alternative explanations (finding two or more different combinations of conditions that are sufficient for an outcome) based in different contexts. An example of such an explanatory theory might go as follows: first, that PB will fail to deepen democracy in the presence of both a very competitive electoral environment and weak civil society where supporters have to spend too much capital developing and consolidating popular support and coalitions; but second, that PB may also fail in a very uncompetitive electoral environment where government has civil society support and therefore the incentive for accountability through participatory democracy for elected leaders is absent. A QCA approach will not see such explanations as 'competing'.

[5] Chapters 6 and 7 walk the reader through this process in more detail when examining the evidence.

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

These assumptions are not the same as those used in detecting interaction effects in regression analyses (see Schneider and Wagemann, 2012: 86–7). Using the QCA approach means we are less interested in causal effects that aim at modelling a meaningful relationship of a single cause with the outcome that is independent of the presence and absence of the other causal conditions (Ragin, 1987, 2000: 95). Rather, causes in different contexts are seen to have effects that do not 'average out' in a very meaningful way. There are two contrasting basic mathematical approaches to these logics that we can distinguish, one formed on the operators of Boolean algebra and the other on the operators of elementary algebra (Thiem et al, 2016).

The methodological assumptions may seem abstract, but they have practical consequences – not only for theories, but also for intervening in the real world if causation is assumed to be neither 'unifinal' nor 'linear'. To put it in more applied terms, I am not looking for answers of the kind: 'controlling for other factors an average 1 per cent increase in political leadership's commitment to participatory methods will lead to an average 6.54 per cent increase in citizen control of collective decisions'. Although this might be a delicious result for current fashions in journal publications if it was obtained through reasonable proxies for the concepts, the practical advice that can be construed from such a result alone is questionable. I am not hoping to subsequently drag a bunch of policymakers and activists into a room and tell them to focus their energy on increasing the participatory commitment of leaders because on average, most of them will be marginally successful. This would not be terribly convincing, and is probably a bad idea for many other reasons. Instead, the importance of context decrees that to generate an outcome it will not do to merely always reduce or increase the presence or absence of a single strong, significant explanatory variable. Increasing the interference of political leaders even committed to participation can have negative effects in certain circumstances. Politicians can derail a process that is otherwise effective where it has previously been insulated from some kinds of polarization. There is currently a burgeoning of work on mini-publics that tries to understand why input from politicians seems to work well in some contexts but not in others.

One answer to this problem of asymmetric hypothesis (where it is reasonable to assume the cause may, for example, improve or worsen democracy) is to try and express expectations about particular conditions more probabilistically. That answer is no doubt always useful. Most generalizations we make are, to some extent, implicitly probabilistic. A popular advance to explanation in the social sciences is to be more explicit about the accuracy of the probabilities based on

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

evidence. Another (not mutually exclusive) answer is to sort through alternate potential explanations and better specify the contexts in which potential causes may act in opposing directions in relation to the outcome. The latter is the focus of the QCA approach.

QCA explanations then allow for causal *asymmetry* and also *multifinality* – Rather than the presence of a condition connected with one outcome being taken to mean that its absence must also be connected with the absence of the outcome, the same condition can be understood to contribute to, for example, deepening democracy in one context while contributing to negating democracy in the presence of other conditions. For example, we might think that a combination of strong financial backing and political leaders committed to participation will explain effective PB, but it does not necessarily follow that a weak financial position will negate PB success. Voluntary support from civil society may pick up the slack where government capacity is curtailed. And a strong financial position may actually weaken PB where leaders are not committed to participation and can freely spend funds on other projects. These are examples of hypotheses I test in the final section of the book.

Intersection, union and negation

Operations based on *intersection* and *union,* as well as *negation* of sets, can help us logically interpret how different conditions combine with outcomes across cases. Intersection of sets corresponds to the logic of 'AND' statements – the intersection of set A and set B contains only all the elements of set A that are also in set B. Union corresponds to logical 'OR' – the union of set A and set B contains all the elements that are in set A as well as all the elements that are in set B. Negation corresponds to logical 'NOT' – the negation of set A (not A, sometimes written as 1-A) is its logical inverse. Throughout I adopt the common notation where the tilde '~' preceding a condition denotes absence (negation) of a condition, the asterisk '*' denotes an intersection of sets (conjunction of conditions), '+' denotes a union of sets (disjunction), '←' denotes is necessary for, and '→' denotes is sufficient for.

Logical cases, remainders and truth tables

A further important aim of the QCA approach is to be transparent about the empirical relevance of the data that has been accessed vis-à-vis what data remains beyond the reach of the research. Any increase in the number of conditions examined in a single analysis leads to an exponential

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

increase in the number of *logical cases*[6] (combinations of conditions that are logically possible; see Ragin, 2008: 24). That is, assuming our number of cases is not very high, if we try to throw more potential causes at the model we will have to admit there are lots of possible combinations of conditions we have no information about. With a small-to-medium number of cases, the number of *logical remainders* (logical cases for which we do not have empirical examples) will usually begin to outweigh the number of potential combinations of conditions for which we have empirical evidence as soon as there are more than three or four conditions in the model. Therefore, a robust medium-N QCA requires the theoretical skills to select and define three to five key explanatory or influencing conditions to be examined with the outcome condition.

Even with a few conditions and a medium-N of cases, the number of potential combinations of conditions runs too high for most researchers to analyse without good visual aids. Here, the *truth table* proves an extremely useful analytic tool. Introduced at greater length in Chapter 7, the truth table concisely maps data to concepts of interest. It summarizes the logical combinations of presence or absence of conditions in each of its rows. The truth table provides information on which empirical cases best ascribe to each logical case, including displaying the 'empty' rows that signify logical remainders. It also illustrates the extent to which cases that are logically similar in their conditions consistently lead to the outcome of interest. The truth table allows for transparent dialogue between theory and evidence.

Fuzzy sets, membership scores and calibration

It is unusual in the social sciences to find that all cases under investigation can be neatly indexed to fall into discrete categories. While good social theory establishes relatively clear conceptual boundaries, we also know that it is possible to identify important differences in degrees to which some condition can be deemed to be present or absent in a case. *Fuzzy logic* allows set-theoretic analysis to account for differences in degree of presence or absence of a condition within a case. Fuzzy sets are further illuminated with the use of examples in the next chapter, but it suffices to explain here that fuzzy sets enhance QCA by allowing for cases to have different degrees of *membership* in a set representing a condition's presence or absence in the case. Where a condition is observed to be present in a

[6] For each k number of conditions tested in any QCA model there will be 2^k logical cases.

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

case, the case can be said to have a full membership (membership score of 1) in the set that represents that condition. The Pope could be said to have a set membership score of 1 in the set of Roman Catholics. Where a condition is completely absent from a case, the case can be said to have full non-membership in the condition (set membership score of 0). The Queen of England could have a set membership score of 0 in the set of Roman Catholics. Other cases may have partial membership scores between 0 and 1, which better describe their degrees of membership of a set (perhaps reflecting a lapsed or 'à la carte' Catholic).

Conditions of interest are operationalized as meaningful sets by the researcher through a process of *calibration*. Calibration is the scientific process of standardizing instruments for measurement. In QCA calibration refers to the specification of the set, and in particular, the description of qualitative states that correspond to specific membership scores to reflect the different quantities of the underlying concept or condition that might be observed. Case selection is inextricably intertwined with condition selection and calibration of sets. The method of coding a variable known as calibration of conditions is that which distinguishes QCA approaches from other forms of data processing.

Consistency or inclusion, coverage, relevance and proportional reduction in inconsistency scores

Finally, for this whistle-stop tour introducing QCA's main features, it is useful to briefly introduce QCA's further measures for dealing with uncertainty in set relations. In social sciences we are not running experiments in vacuums, and finding a perfect relationship represents a very high standard. There are likely to be some inconsistencies in evidence for a relationship across a number of cases. A social scientist would normally be willing to accept some minor inconsistencies to avoid disconfirming a relationship, especially when armed with a transparent measure of fit (being able to articulate the extent of uncertainty). QCA methodologists have developed several indicators of the degree to which set relationships are met. We can provide a measure of the extent to which data confirm the subset relation with an outcome, and this is known as consistency[7] or inclusion.

[7] It is important to briefly introduce consistency to the reader here to aid understanding of the output. This explanation draws on Schneider and Wagemann (2012, Chapter 5), and a much more detailed explanation and discussion of the concept can be read there.

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

Figure 4.1: Shows two tests for necessary conditions

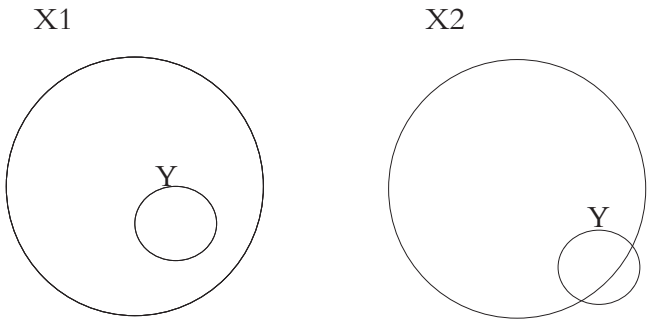

Consistency or *inclusion* is perhaps the most important of QCA measures of fit, and is a measure of the extent to which a subset–superset relation is confirmed across cases. Messy social science data will rarely produce clear outcomes where all cases of a certain configuration behave in exactly similar ways. For analysts, the measure provides a sense of how consistent the evidence is for a certain relation of necessity or sufficiency.

To give an example, the Venn diagrams in Figure 4.1 show two tests for necessary conditions. We can treat each of the two Venn diagrams as separate tests for necessary conditions for Y. On the left, the condition X1 is a consistent superset of Y. Whenever we see Y, we see X1. On the right, we see Y is almost but not fully covered by X2, and is not a perfect subset. Sometimes there is Y without X2. There is some inconsistency in its inclusion within the set. A social scientist might, however, be happy to say that X2 is 'almost always necessary for Y'.

Inclusion scores shown in the analysis are always calculated across all cases. The standard-used minimum consistency thresholds within the QCA literature range from 0.9–0.95 for necessary conditions and 0.75–0.9 for sufficient conditions. Thresholds of inconsistency effectively serve to avoid dismissing superset–subset relations where we have 'very few near misses' across cases (Ragin, 2009: 108). However, appropriate thresholds will depend on the nature of the inquiry (theory-testing work may require a higher threshold than exploratory work), the number of cases and the nature of the data.

Coverage and *relevance* of necessity are both measures of the triviality of necessary conditions (Schneider and Wagemann, 2012: 147; see also Goertz, 2006). Necessity coverage is an expression of how much smaller the outcome set is in relation to the conditions set (how much of an outcome set Y covers a condition set X). Low necessity coverage can indicate trivialness of necessary conditions (X is quite big relative to Y)

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

and high coverage can indicate a relevant necessary condition, provided X is not too close to a constant. *Relevance of necessity* was later developed as a measure of triviality that considers the latter problem caused by a relatively large X set (see Schneider and Wagemann, 2012: 232–7). The *proportional reduction in inconsistency* (PRI) measure is useful for fuzzy analysis of sufficient conditions because it identifies how much a combination of conditions is specifically a subset of the outcome vis-à-vis the absence of the outcome. The aim is to help researchers avoid illogical interpretations of superset–subset relationships.

While I have introduced some of the key concepts and nomenclature in this section, I continue to familiarize the reader with the approach by showing how it can be applied to existing studies of PB in a cumulative way.

Applying qualitative comparative analysis to participatory budgeting

When I began studying PB I noted several commentators asking for systematic comparisons. My attention was drawn to a report (Pratchett et al, 2009) that used a short-form QCA to systematically review 19 PB processes. Pratchett et al offer an early attempt using more rudimentary applications of QCA available at the time to try and uncover patterns of causation in PB outcomes from then available case materials. While the limitations of the time meant some results were difficult to interpret, the report showed the promise of the approach. As QCA's sophistication rapidly increased, I committed to applying it to seek to improve the knowledge base on PB.

The rapid diffusion of democratic innovations, especially PB, and the rapid increase and refinement in QCA techniques both provided me with moving targets during this ongoing research project. As I gathered data, I began to apply QCA to small samples of contrasting cases of PB across different continents, and gradually increased my dataset of cases using criteria that I expand on throughout this chapter. Over the course of this study iteration between analysis, population refinement and conceptualization often led to refinements (see, for instance, an early comparison of six cases in Ryan and Smith, 2012, and 19 cases in Ryan, 2014). Early efforts allowed me to test the effectiveness of the approach, which showed that it was worth persevering with. It also allowed me some time to reflect on population definition, calibration of conditions and presentation and interpretation of outputs. In advance of the software-aided analytic moments of QCA, if we take it to be an appropriate method for a 'medium-range' social science

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

programme of research on PB, we are faced with a number of questions. First, what counts as a case of PB? Second, can we adequately define conditions of interest – both influences and outcomes? And third (and more practically), is the case material or data available suitable to provide enough insight to describe these conditions, both qualitatively and quantitatively?

Iterating between cases and theory to accumulate participatory budgeting knowledge

An advantage of case focus in methodology is to allow serious thought to problems of boundaries of cases (Ragin and Becker, 1992; Byrne and Ragin, 2009). Case-based methods vary in the tools they employ and the number of cases they compare, but they hold in common an understanding of social research that is sensitive to the complexity of the case and the scope of generalizations. To allow for generalizations within a designated scope QCA requires sound theoretical specification of the population of cases and selection of key conditions for comparison. This means that case-based methods are characterized by constant iteration between inductive and deductive strategies,[8] and between evidence and theory. The approach favours ongoing problematization of the scope of populations, causes and outcomes in a piece of comparative research (Ragin, 2000; Rihoux and Lobe, 2009).

What is an appropriate way to demarcate the cases of PB we should include in analysis from those we should not? How do we recognize the phenomenon we are interested in explaining and perhaps predicting? There can be pressure to consider a relatively homogenous population, especially where observations are limited. After all, 'universal propositions may be safely applied to an actual subject matter only in so far as we are thoroughly familiar with the *type* of *object* of which the actual case is a *sample*' (Cohen and Nagel, 1934: 280–1; original emphasis). For the purposes of this research, this involves defining what a case of PB is (and what it is not). A number of challenges present themselves. PB is still a relatively young concept (compared to, say, election, parliament or protest). How exactly does it successfully distinguish itself from other traditional political processes or democratic innovations and ongoing methods of participation and/or governance? A variety of definitions of PB are to be found in the literature. Many

[8] I take induction and deduction here as common placeholders for strategies towards knowledge accumulation that prioritize hypotheses generated from empirical investigation and from theory respectively.

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

definitions list several features, for example, the facilitation of regular meetings, the provision of information, direct participation, self-regulation and so on (see Chapter 2). Definitions will often suggest the presence of all features in combination is essential for a case to be distinguished as PB. Many observers either explicitly or implicitly use a much lower number of conditions, or take the inclusive line that any participation in budget matters by 'ordinary citizens' constitutes PB.

A stricter approach is perhaps favoured by area scholars and those inspired by the design of the Porto Alegre case. It may also be difficult for researchers, many of whom have normative commitments to deepening democracy, to concede any ground on the attractiveness of what 'participatory budgeting' was before the concept migrated and developed, or perhaps regressed. Nevertheless, the position of the Porto Alegre case as not just a poster case, but moreover, the archetype of PB, is a challenge in many aspects for a comparative research agenda. What value is there in defining the population based on the best or first case? Polymorphism is a feature of diffusion (Porto de Oliveira, 2017: 72). PB has necessarily been implemented in different ways, as it has been adapted to different contexts. Processes that are called PB in one country would be called something else in another (Sintomer at al, 2013: 2). Röcke holds that for PB, '… the transfer or diffusion of institutions, ideas or practices always necessitates at least a minimal amount of re-interpretation with regard to personal strategies and characteristics of the new context' (2014: 28). 'Pragmatic experimentalism' across cases is also an important feature within Brazil (Baiocchi et al, 2011: 81). The question is whether this reinterpretation should represent variation in measurement, or whether there is such a variation in quality that the phenomenon is no longer recognizable.

The success of PB has led to political opportunism and promotion of the concept, but not necessarily any single model. It is difficult to know whether the adaptation of PB in new locations is a case of well thought out revision of a concept incorporating local knowledge, or the muddled product of garbled worldwide translations. It can be equally foolish to define cases based on minimal criteria. A plethora of programmes now called PB may be PB in name only, based on the fame of the original case and the allure of being seen to implement an 'innovation' (Ryan and Smith, 2012). This discussion highlights the challenge within such a comparative study of distinguishing variations in quality from variations in kinds of phenomena. The question is often asked as to whether structured comparison is a good approach given the emerging nature of the phenomenon. The answer is that structured comparison is always a good approach. If the approach can

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

help identify and separate differences in kind and in quality, it should provide a development on previous efforts at cross-national comparison.

Scoping the population of cases and conditions of interest is never an easy task in emerging subfields of research. Yet, by its nature, a newly emerging phenomenon will be limited in its scope for comparison to a small to medium-N of cases. This is certainly true of participatory democratic innovations such as PB when compared across countries (Ryan and Smith, 2012). The field is (over)loaded with good case studies and small-N comparisons that rely on in-depth ethnography-based methodologies. This has led to a call for more systematic cross-case comparisons (Smith, 2009; Spada and Ryan, 2017). In the previous chapter I discussed how the rapid diffusion of PB has led to increased problematization and contestation of its core conditions. How, then, can a researcher adequately scope their population of cases for a valuable systematic comparison enabling modest generalization?

One obvious strategy employed in small-N case-comparative research is to limit cases by keeping key variables constant to control extraneous variance. This logic is often applied in the area studies approach or by limiting comparative designs within one country or system of government (Lijphart, 1971: 688). Yet small-N research based on researchers' in-depth case studies is often criticized because of its alleged increased risk to subjective bias. Moreover, cumulation of small-N research for comparison across studies can be difficult because of researchers' particular and differing intentions. In Peters' words sometimes it is argued that 'the researcher is the major independent variable' (2013: 169). The research I present here relies on my interpretation of the work of others. This is an advantage, because I can cumulate, compare, triangulate and cross-check. It might also be seen as the worst of both worlds – one person interpreting the subjective biases of others to suit their own. Transparency is the researcher's weapon to fight that impression, and I hope that I offer value in that respect.

The challenge of expanding the evidence base

Studies of PB have been either narrow and deep, or wide but shallow in their treatment of the cases. Up to now, description and categorization of several cases has relied mostly on induction, with several compilations that bring together descriptions of similarities and differences among cases. This is partly explained by the way in which puzzles present themselves to researchers and partly by the tools we have available. In contrasting cases across Asia and Europe, Sintomer et al bemoan that there is no 'uniform model to which the others could be compared'

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

(2013: 20). They offer some organizing principles – their comparison of European cases offers the closest application yet to some middle path. Working back from ideal types they offer a PB typology categorized according to different features of origin, organization, deliberation and the role of civil society (2016). Beginning from those ideals, relative distances from the ideal types can be roughly approximated and conveniently visualized. This approach does not allow for systematic or transparent measurement, and is difficult to replicate. Cases tend to cluster near types according to their national context, so there is a danger that typecasting of cases according to their national backgrounds occurs using such an approach. To my knowledge, no one has come close to expanding on a full typological theory based on evaluation of a property space in the manner suggested by George and Bennett (2005: 257).

There is also a familiar issue of accumulation of research knowledge that is biased by global power structures within the academy and research and governance professions worldwide. Despite excellent efforts in sharing of information on several interesting examples of PB in East Asia, much of the work going on in the region remains largely invisible to scholars of PB (Cabannes, 2018: 243). PB represents a 'return of the caravels' and much international exchange is done in Lusophone, Francophone and Hispanic spheres. Nevertheless, in-depth casework that reaches a global scholarly audience is always biased toward spheres where Indo-European languages dominate, and in many countries scholarly work is still rewarded by a publishing regime where English is the lingua franca. Peters warns that we should be wary that what are seen as popular generalizations are not simply artefacts of the tendency for Americans or Britons who still dominate elements of the discipline to select cases for ease of language use (2013: 171). It is one reason why we require and value the work of skilled area scholars.

Area research is a double-edged sword. Where comparisons have been more systematic, they have tended to select cases within continental bounds. Despite unquestionably high-quality contributions, most comparative work on PB has found ways to find solace in the unique and unusual aspects of the areas they control for. This can come in the form of stressing the uniqueness of the Brazilian political system and culture (Wampler and Avritzer), focusing on national framing (Röcke) or employing approaches that are more likely to uncover differences than uncover similarities across continents (Sintomer et al). The field can benefit from making fewer assumptions about the differences among PB that take place in different macro-comparative units of political organization (or at least testing these

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

assumptions systematically across cases). To be clear on what is at stake in defining PB for comparative research – and this is very important if actors capable of having political impact are able to understand the scope of research findings – we must systematically compare the key elements of phenomena in order to define the concept. People get away with calling any old consultative event a 'successful' democratic innovation partly due to the absence of general systematic analysis and descriptive conceptual work to set things straight. By systematically comparing programmes, we can show that 'anything goes' has very different necessary and sufficient conditions to empowered PB. These differences can be observed and described across several cases and contexts, and by doing so we provide a clear vocabulary to those wishing to expose flawed participatory efforts. Only by seriously attempting systematic cross-regional comparisons will greater insight be gained into what specificity of regions and countries is lost in abstraction for comparability. We can also challenge the assumption that processes in one part of the world are different to another (at least in terms of their most important characteristics that significantly influence empowered participation).

The process of conducting QCA involves both construction and deduction, allowing constant reflexive iteration between theoretical assumptions and measurements. The method has an inimitable way of constituting clarity about theoretical assumptions involved in measurement, selections of cases, choices between parsimony and complexity in description of empirical regularities, and interpretations of results (Rihoux and Lobe, 2009: 237). Much of the procedure I illustrate will seem familiar to practised small-N researchers. This kind of careful iterative reasoning is their stock and trade. So, what's the big deal? Applying the logic of case-based research beyond conventional boundaries in the number of cases that can be compared in any one sitting can lead to unique insights. The value of a medium-range social science is that it retains the transparent conscientious reflexivity of the case researcher (Flyvbjerg, 2006: 235), and applies this contemporaneously with the expansive lessons derived from testing relationships across several accumulated instances of a phenomenon.

In case-based research, the addition of new cases one by one is challenging. This is intensified when dealing with a larger-N. Arriving on new cases may have consequences for both population definition and variable definition. New cases may present a variation in the unit of analysis significant enough to force the researcher to reconsider the population. Several new cases of PB that appear substantially different

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

may lead us to establish new classes, concepts or types (as seen in the work of Sintomer et al and Cabannes and Lipietz, outlined previously). Another way of understanding the effect of discovery and comparison is that new cases may require re-evaluation of distinctions between scope conditions and influencing conditions (see Walker and Cohen, 1985). We might find we have to adapt our theories to distinguish certain conditions as influencing an outcome within the scope of others – finding one constellation of conditions helps to establish successful citizen participation in one context, and another elsewhere. For example, our hypothesis might be that political and financial support are necessary in poorer contexts, whereas civil society demand is an essential explanatory component in combination with political support in richer municipalities. My approach facilitates that process of judging alternative hypotheses by identifying logically sufficient combinations of conditions across contexts.

On the other hand, the understanding of how conditions are observed in new contexts may force the researcher to re-evaluate the operationalization of variables – how they have calibrated their conditions. While this iteration between cases, measurement and theory is time-intensive, its transparency in construction of the research is a key methodological advantage of QCA. I do not wish to claim that iterating between measurement, concepts and samples is peculiar to the approach. Some proponents of QCA have implied this at times, and it is unhelpful because all research involves elements of those procedures. It is better to say that the QCA approach has some heightened qualities in its sensitivity to the problematization of these issues. It is also worth pointing out that this problematization can lead to over-conscious thinking (on the latter problem, see Sartori, 1970).

Some definition is an important starting point in conceiving of a universe of cases. We saw earlier that definitions of PB tend to round on a list of key fundamental characteristics. Those provided by Wampler and Avritzer evoke some of the key design features that have contributed to distinguishing PB as an innovation in participatory governance. They are slightly abstract and tend to favour the conditions of Brazilian designs. Not surprisingly the best starting point for worldwide comparison comes with Sintomer et al's definitions listed earlier. These conditions present sensible criteria for distinguishing PB from more common 'traditional' consultations with citizens outside of the political classes.

The final sections of this chapter explain the logic in scoping and selecting cases for the analysis of PB presented in Chapters 6 to 9 of Part III.

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

Differentiating PB from participatory grant-making processes: what PB is not

Sintomer et al's criteria would work to exclude from the population many of the cases named 'participatory budgeting' in England (Hall and Röcke, 2013; Ryan, 2014), and similar examples elsewhere, which tend to involve participation in the distribution of ad hoc small grants. These processes also tend to involve administrative jurisdictions that are not governed by a specific representative political body (for example, neighbourhoods, highways).

The English cases are instructive for conceptual clarity. To give an illustration, a suburb in the city of Southampton (where I live) held a 'PB' process that started in 2008 and continued well into the next decade. Thornhill, with a population of about 10,000 residents, was identified in 1999 as one of 39 exceptionally deprived areas of the UK as part of the then Labour government's major regeneration programme called New Deal for Communities (NDC). Each area was granted approximately £50 million worth of government investment, over a 10-year period (Fordham, 2010: 12). As part of a deal struck with the board controlling NDC funding allocation, the local primary care trust (PCT) agreed to make a recurring financial contribution to the area distributed via a PB process (Bonaduce de Nigris, 2008: 4). Support for the process was also provided by the city council. The process involved local residents creating projects and bidding for money from a single pot of what at its height was tens of thousands of pounds. These projects were then presented at a public meeting in the local school, and residents voted on which projects they would like to receive funding. Their votes were tallied, and project funding allocated accordingly.

This is typical of many small projects inspired by PB in the UK (PB Unit, 2008), and is similar to many described in other parts of the world (see, for example, Shah, 2007). I do not wish to suggest that these processes do not have value. From my observations, the process in Thornhill incentivized creative community-led projects, improved community cohesion, gave some decision-making control to residents, and gave participants some sense of the trade-offs involved in making decisions to allocate scarce resources. Shorter pilot programmes also mean shorter time frames for demonstration effects, which could, in theory, lead to pressure for greater empowerment – although evidence for this mechanism appears rare.

When we try to compare this process with PB in Porto Alegre, we are comparing apples with apple seeds. There is no hope for participatory grant-making processes like these to achieve empowered participatory

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

democracy without them being planted more firmly within the legitimate democratic architecture. The process in Thornhill fails to reach at least four of Sintomer et al's criteria. It is based on a small suburb of roughly 10,000 inhabitants, no more than a portion of a local city council ward. The amount of money was set by funding schemes that ordinary citizens had no control over. Projects were limited to those aimed at improving health and wellbeing, so the agenda was set narrowly elsewhere (by the constraints of the funding stream). There was no formal public deliberation – in fact, officers felt that 'Q&A sessions would be difficult to manage, especially as contentious questions might be asked and issues of fairness might arise' (Bonaduce de Negris, 2008: 17). There was only one public meeting a year, and the continuation of the programme was completely dependent on the goodwill and fiscal priorities of funders. Those in the UK who were promoting a programmatic approach hoped that these small projects were pilots that would lead on to PB proper, but the projects' logic resembles more the 'do-ocracy' popular in the Low Countries than meaningful power-sharing between state and citizens. The prospects for programmatic 'mainstreaming' in the UK have now definitely moved north of the border into Scotland (see Escobar et al, 2018).

Other processes in England come closer to satisfying the criteria but still fall short. For example, in the London Borough of Tower Hamlets in 2009 and 2010, a 'PB' process allocated £5.04 million of council services by asking local residents to come together, deliberate and vote on priorities in their neighbourhoods (U-Decide, 2009). The Tower Hamlets process was not a grant-giving process and involved residents voting for money for their preferred service. Despite constraining its agenda to bids for 'additional' public services, it opens up a far more significant amount of public funds to local decisions than the grant-making process described earlier. However, the institutional design of this process suggests it is better described as a 21st century town meeting with a focus on local budgeting. Meetings were not regular, and there was no opportunity for the kind of civil society investment over a year-long period expected in PB. The process did not allow in its structure the development of accountability across budget delegates at neighbourhood, region or city or borough level. It was discontinued after 2010 on grounds that such funds were no longer available in a meaningful way as a result of the financial downturn. Although Tower Hamlets PB spend was drawn from a 'mainstream' budget and not based on the ad hoc 'funny-money' of one-off grants or surpluses, as is often the case in UK PBs (see Hall and Röcke, 2013: 195), its foundation was still very insecure.

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

This vignette invokes another important criterion for definition of the PB universe I use for comparison. Those researching democratic innovations have struggled to explain their populations by differentiating failed cases from non-cases, which partly explains a skew in publications emphasizing successes (Spada and Ryan, 2017). Like many PB scholars, I am interested in explaining how citizen control becomes institutionalized in a PB process. In the case of Tower Hamlets, a key reason why control could not become institutionalized was that the programme was discontinued after two years. The two-year criterion is one I use to differentiate a negative case from a non-case of PB. I assume that a case that is not institutionalized to recur more than once is not useful to explain how citizen control of PB processes happens. These early abandoned cases would certainly be useful for explaining why PB is initiated, and why it survives or does not. I would also not deny that one-off events can have many good democratic outcomes. My argument is that we cannot understand those processes as PB. If I had unlimited resources, I could go looking for lots of material to test the assumption that there is a complete affinity between recurrence and PB-style empowerment. I assume PB is a process involving a longer than two-year commitment with regular meetings. Processes that do not recur have very different qualities that I am not able to explore at length here.

Cases and data for a worldwide comparison

Ever more sophisticated case selection has helped correct early narratives in democratic innovation research that relied on best-case examples. Case selection (sampling) requires judgement, and knowledge of cases and their relationship to the population will provide the grounds on whether more or less modest inferences can be made. Early monographs on QCA were keen to stress that populations could never be taken as 'given', where given meant that populations were not expected to change during the research process (Ragin, 2000: 45). The basic point Ragin wanted to make is that this process of 'casing' and determining populations is itself an opportunity for theoretical advancement (2000: 54), as cases are discarded, added, revised and refined throughout the research. One advantage of QCA as an approach is that it did not seek to artificially separate or collapse comparison as a conceptual exercise from comparison as causal analysis. That is part of its attractiveness for assessing PB.

Nevertheless, many applications of QCA have had a disappointing amount to say about case selection. When I first broached experienced

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

QCA practitioners with my sampling questions early on in my research, their unspoken assumption seemed to be that QCA was a home for research that involved analysis of 'full' populations that might be medium-N (for example, Western European welfare states). There remains less work, to my knowledge, that leverages transparency in information-oriented (deliberately non-random) sampling from a much larger population. There has also been a recent trend to apply QCA to so-called 'given' populations where sampling is not really practical or desirable, or to apply QCA to large-N samples. Those analyses aim to leverage QCA's set-theoretic analytic component, often producing quite useful results. The latter trend has led some to write sensible summaries distinguishing differences in QCA as both an approach and a set of techniques (Schneider and Wagemann, 2012; Thomann, 2020; Thomann and Maggetti, 2020).

Sampling is inherent in every piece of research and samples derive from a combination of population definition and case selection. I am transparent about the practical considerations that influence construction of a model to explain PB outcomes, and provide the reader with the relevant information to evaluate my inferences and to contribute their own. The decision to include or exclude cases is important because anyone evaluating the model should be aware of how sensitive model 'solutions' are to evaluate inferences based on those decisions. Therefore, I explain, particularly in Part III of this book, how cases that are not included contribute to, or might modify, general statements. The *truth table* is a useful tool here, whereas the computerized *minimization* in QCA (where an algorithm removes irrelevant terms from explanation) can be more obfuscatory if not well explained. I explain how these tools can be used in an open manner.

I set about building a dataset with enough cases to make meaningful comparisons beyond a small-N, incorporating substantial variety. I also wanted to apply fuzzy set QCA, which meant I needed to be confident that I could obtain a requisite quality of data on important conditions I wanted to test and explain for each case. When I first approached this project with some knowledge of the rapidly emerging literature, I had a sense that the field was beginning to ripen for such a task. As PB began in Latin America and has taken hold in all corners of the globe, much of the literature appears in a variety of languages. I was limited mostly to English language sources (through my own linguistic ignorance), but the research was undertaken with the premise that this 'ripening' was occurring as a certain quantity and quality of knowledge on PB was becoming available across English language sources.

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

For an innovation that is not a mere technical fix or upgrade, the diffusion and increase in the number of cases of PB across the globe has been staggering. Porto de Oliveira (2017) provides an impressive narration of the threads of diffusion through international relations among cities, and the pivotal role played by actors – politicians, officials, scholars and activists. It is close to impossible for even skilled, resourced and well-coordinated researchers to keep track of all the cases of democratic innovations as they diffuse and diversify throughout the world (Smith, 2009). The classification 'innovation' is itself to some extent unhealthy in that regard (Smith, 2019). Classifying something as an 'innovation' by definition requires ignorance on the part of many toward it, and nowadays almost always connotes assumptions that its solutions to problems are a fait accompli (Spada and Ryan, 2017). Inadequate conceptualization of cases of participatory governance has been blamed on some pathologies in research design in participatory governance (Spada and Ryan, 2017; Ryan, 2019). Despite those difficulties, some efforts in data collection have been impressive, such as that of Participedia (participedia.net),[9] which crowdsources information on a variety of innovations providing sometimes quite high levels of information on cases. While PB can still be considered an innovation in several contexts where it is rare, it is certainly one that is now well into stages of adoption. Another impressive data collection effort is the 2019 PB Atlas, which relies on an expert survey to estimate the number of participatory budgets worldwide (Dias et al, 2019). The collaborators estimate over 11,690 cases. That figure is based on a more inclusive definition than I provide, and comes with large confidence intervals, but nevertheless shows the breadth of influence and impressive spread of the idea.

Despite these advances, until now a medium-range study including both breadth and depth had been difficult to imagine. Retaining the value of case knowledge and comparing across more than a handful of cases is time-intensive. For some scholars of participatory governance, combining quality and quantity is a life's commitment and usually involves some research help (see Bryan, 2004, for an example of such an approach in his study of town meetings in New England in the US). In order to build a dataset, I repeatedly scoured online sources at regular intervals and gradually built up contacts across networks to access new information, recording sources and data as I went along.

[9] I have worked with the Participedia project since 2009.

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

In the wider-ranging literature on worldwide cases I have mentioned many cases appear as cameos more so than in-depth case studies. Therefore, as I collected data on mounting claims of potential conditions that purportedly explained many important outcomes, the empty cells in my dataset were increasing exponentially rather than decreasing. I began to narrow my focus on what cases and conditions appeared most likely to bear enough information for qualitative comparison. In some literature, even a short article would provide very clear descriptions of conditions I was interested in. At other times researchers documented quirks of particular cases that were less useful for cross-cutting research questions, and did not describe a case's basic elements. In some of those cases I was able to find what I needed through cross-checking sources or clarifying an interpretation through communication with the field researchers. The 30[10] cases selected for the analysis presented in this book represent some of the sheer variety of cases and sources that any comparative study of this type must work with. Although I am limited in the main to English language sources, the cases are not biased towards the English-speaking world. The data contains 13 Brazilian cases where the innovation was strongest in its early waves and sources are several, and 6 in Peru where the first national law was put in place, and Stephanie McNulty's work provides some clear relevant evidence on those cases. These are compared with 2 cases in France, 2 in Spain, 2 in the US, and 1 each from Argentina, Canada, Italy, Germany and Uruguay. The cases and some of the main sources from which information was garnered are listed in Table 4.1. Contributions ranged from books treating a single case (for example, Abers, 2000), books treating a small number of cases in comparison (for example, Talpin, 2011; Röcke, 2014), to some mere snapshots, but these 30 cases represent those for which there was enough information of enough quality to make valid inferences about the conditions that were necessary and sufficient for citizen control of collective decisions. The list is far from exhaustive, but includes sources from which the analysis I present here could be performed to a minimally acceptable level of validity.[11]

[10] In the main analysis presented in the following chapters, the Recife case is split into two cases representing the pre- and post-PT governance eras. The approach is taken by Montambeault (2016), and its justification is discussed in the following section. Therefore I analyse 31 cases in total.

[11] I also had formal research interviews about cases with Josh Lerner, Anja Röcke, Dennis Rodgers, Julien Talpin and Brian Wampler. I tried to stay as true as I could to the descriptions in peer-reviewed materials. All errors of interpretation of material are, of course, mine and mine alone.

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

Table 4.1: Cases and main sources for the analysis presented in this book

Case	**Main source(s)**
Ayacucho 2003–06	McNulty 2011
Belo Horizonte 1993–2011	Nylen 2003a, Wampler 2007a, Avritzer 2009, Montambeault 2016
Berlin-Lichtenberg 2005–10	Röcke 2014, Herzberg 2013
Betim 1993–2000	Nylen 2003a, 2003b
Blumenau 1997–2004	Wampler 2007a
Buenos Aires 2002–09	Peruzzotti 2009, Rodgers 2010
Cajamarca 2003–06	McNulty 2011
Camaragibe 1997–2000	Baiocchi et al 2011
Chicago 49th Ward 2009–13	Baiocchi and Ganuza 2016, Pin 2017, Weber et al 2015
Córdoba 2001–07	Baiocchi and Ganuza 2016
Cusco 2003–06	McNulty 2011
Gravataí 1997–2000	Baiocchi et al 2011
Ipatinga 1998–2004	Wampler 2007a
João Monlevade 1997–2000	Baiocchi et al 2011
Lambayeque 2003–06	McNulty 2011
Loreto 2003–06	McNulty 2011
Mauá 1997–2000	Baiocchi et al 2011
Montevideo 1990–99	Goldfrank 2011
Morsang-sur-Orge 1998–2006	Talpin, 2011
Moquegua 2003–06	McNulty 2011
New York 2011–16	Gilman 2016, Swaner 2017, Jabola-Carolus 2017, Su 2018
Poitou-Charentes 2004–08	Röcke 2014
Porto Alegre 1988–2004	Abers 2000, Baiocchi 2005, Santos 2005, Gret and Sintomer 2005, Wampler 2007a, Avritzer 2009, Goldfrank 2011
Recife 1993–2000 and 2000–08	Wampler 2007a, Montambeault 2016
Rio Claro 1997–2004	Wampler 2007a
Rome Municipio XI 2003–06	Talpin 2011
Santo André 1997–2001	Wampler 2007a
São Paulo 2001–04	Wampler 2007a, Avritzer 2009
Sevilla 2003–06	Talpin 2011
Toronto Community Housing 2001–10	Lerner and van Wagner 2006, Foroughi 2017

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

Readers familiar with the cases will already have noticed that there is a wide variety of cases in different contexts among the sample. The cases include some well-known municipal participatory budgets, but with different designs in different contexts, one focused on housing tenants, another on high school students, at least seven regional ones, some extremely large cities, and some whose jurisdiction is below the municipal level. Some of the cases are not always called PB in some of the instances where they are presented. All cases fit within the scoping criteria outlined previously.

Comparing diverse cases is an overlooked strength in much comparative research, where an obsession with mimicking control has become prohibitive to some interesting research designs. I specify the time frames covered by the data for each case. Where reasonable, I split cases into separate time frames. This is done only where field researchers suggested a clear distinction within the case. Some of the cases rely on fieldwork that was carried out some time ago, some at the height of the success of PB in Latin America in the early 2000s. Some of the North American material is more recent. Many of the earliest monographs using QCA were comparative historical analyses (Wickham-Crowley, 1991; Brown and Boswell, 1995), and there are elements of that approach in the work presented here. There are clearly different levels of information for each case. Uneven amounts of knowledge relating to each case are an accepted (if not desirable) characteristic of case-based comparative research (Newton, 2006: 851). I make every effort to be clear in analysis and discussion of cases where information was more or less lacking. Many interesting cases have been left out for lack of some authoritative information. In Chapters 8 and 9 I discuss how I imagine these cases would have contributed to the analysis.

Some readers may also question the desirability of including Porto Alegre. As the archetype of PB, it has been seen to have much influence on other cases and models of PB, and is still often seen as a yardstick for many other programmes to aspire to. It may be considered an 'outlier'. In many ways, the success of Porto Alegre is a necessary condition for any of the other cases in the dataset to exist. The question we must keep in mind here is whether this influence is an enduring significant factor affecting outcomes that vary across cases. Using QCA we can try to evaluate to what extent Porto Alegre compares with other cases – whether other cases are similar in key characteristics – and we can then return to the case and examine whether the findings make sense in light of our knowledge about it and other cases. We now turn to the substantive questions we would like answers to across these cases.

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press

https://doi.org/10.46692/9781529209938.007 Published online by Cambridge University Press